

MXNet: Lightweight, Flexible, and Efficient Deep Learning Library

Naiyan Wang

TuSimple

MXNet is developed by over 100 collaborators

Special thanks to

Tianqi Chen

UW

Mu Li

CMU/Amazon

Bing Xu

Turi

Chiyuan Zhang

MIT

Junyuan Xie

UW

Yizhi Liu

MediaV

Tianjun Xiao

Microsoft

Yutian Li

Stanford

Yuan Tang

Uptake

Qian Kou

Indiana University

Hu Shiwen

Shanghai

Chuntao Hong

Microsoft

Min Lin

Qihoo 360

Naiyan Wang

TuSimple

Tong He

Simon Fraser University

Minjie Wang

NYU

Valentin Churavy

OIST



Ali Farhadi

UW/AI2



Carlos Guestrin

UW/Turi



Alexander Smola

CMU/Amazon




Zheng Zhang

NYU Shanghai

Deep Learning

Image Understanding




Speech Recognition



Natural Language Processing



“Deep Learning” trend in the past 10 years




Packages



theano

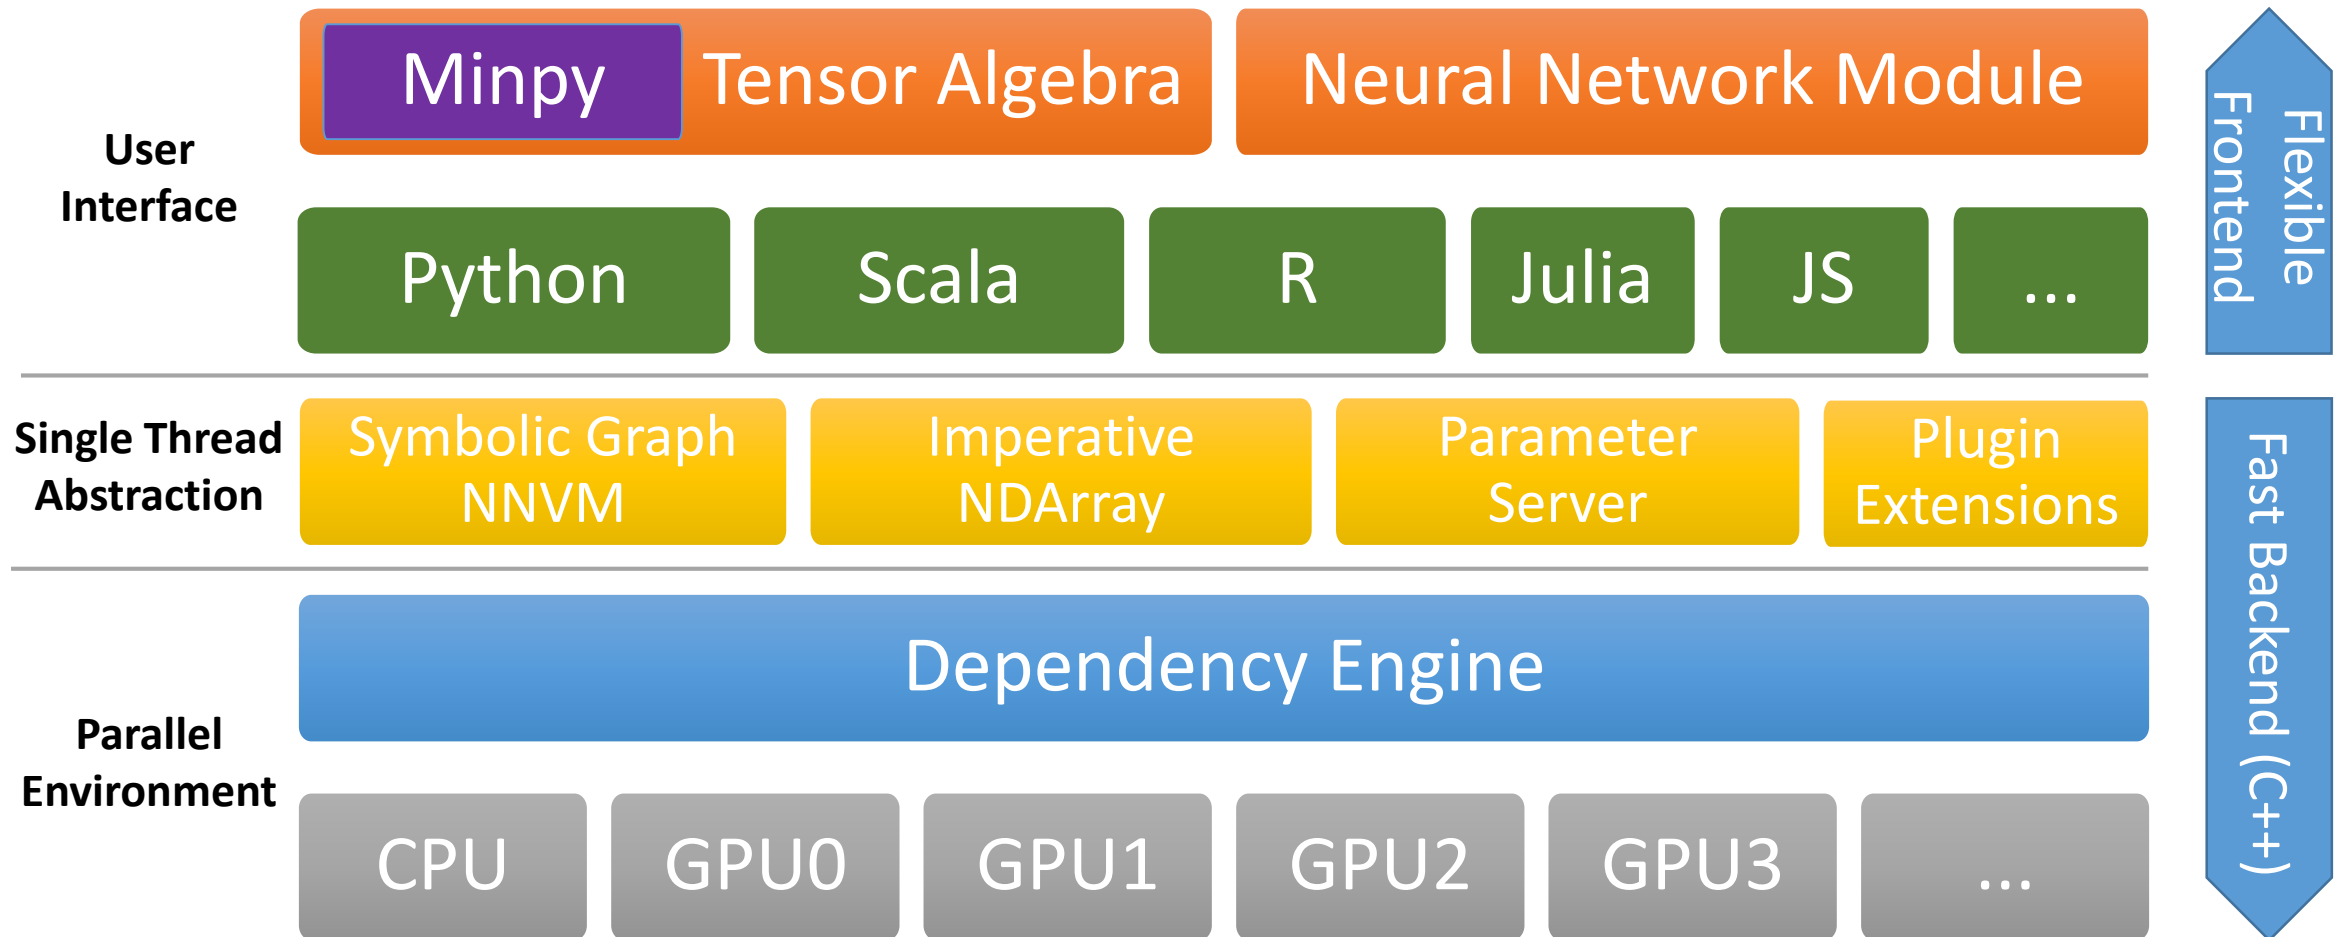
Caffe



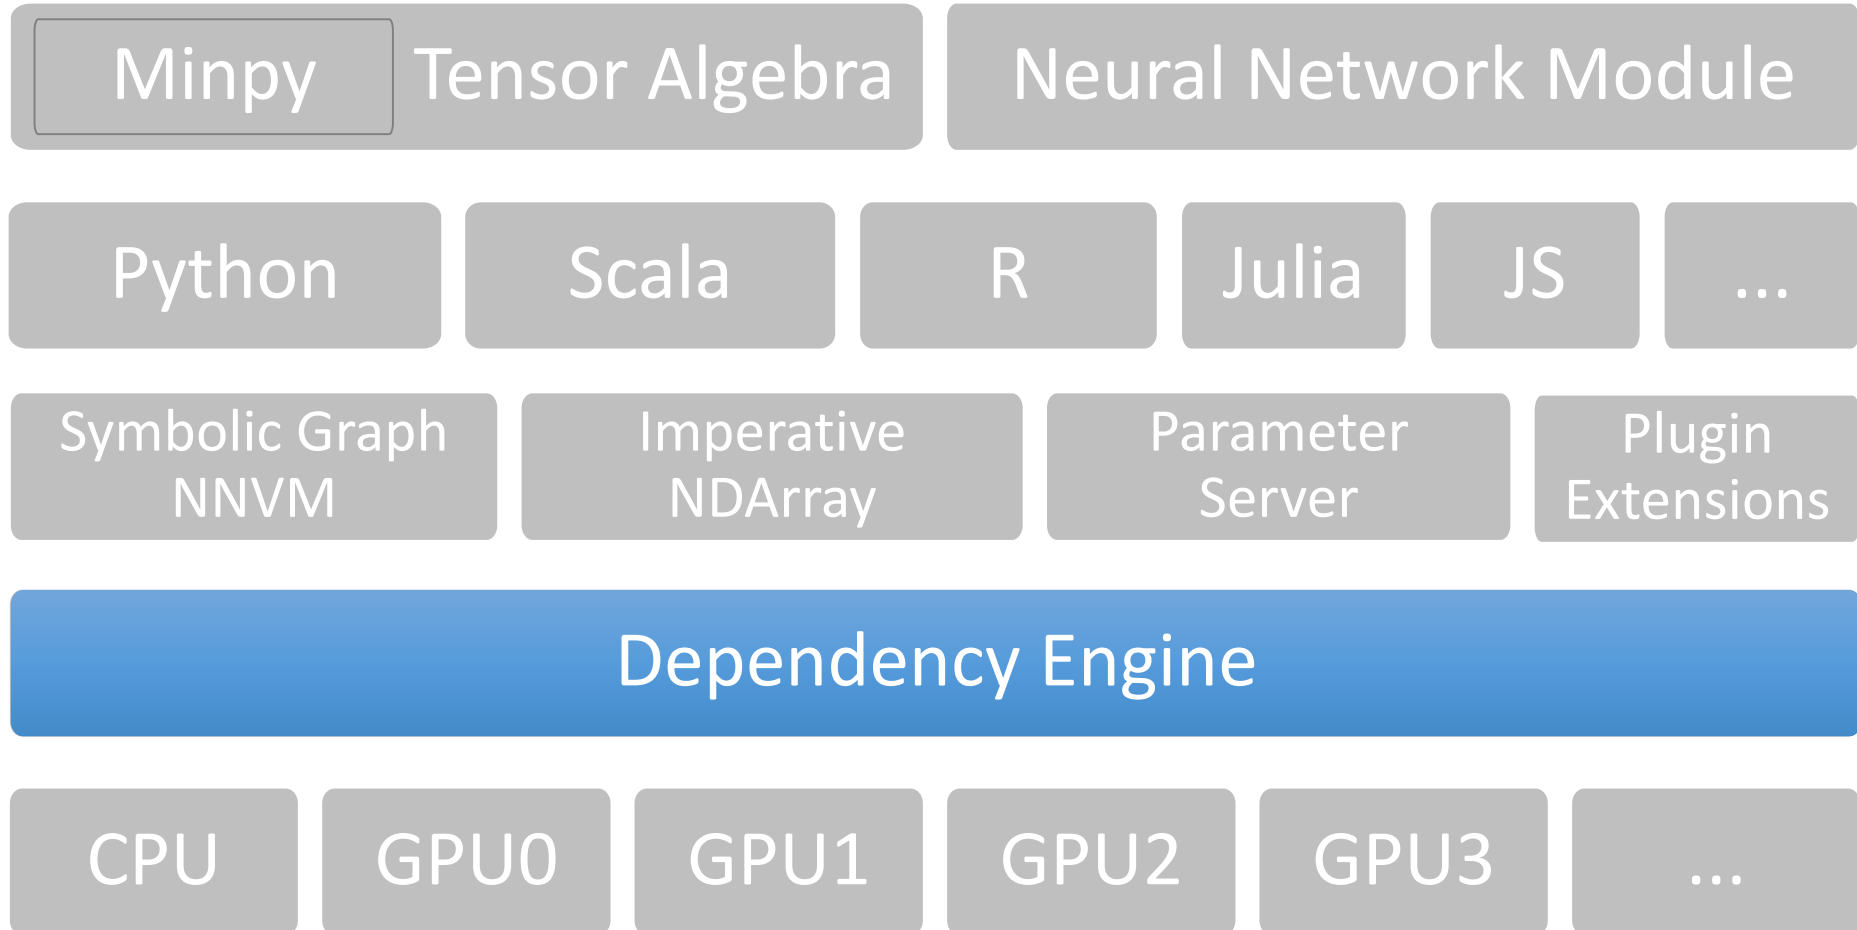
MXNet's Approach

- Lightweight
- User defined flexibility-efficiency trade-off
- Transparent scaling
- Deploy everywhere
- Modular and extendable
- Mixed declarative and imperative programming
- Community driven open source

The MXNet Stack


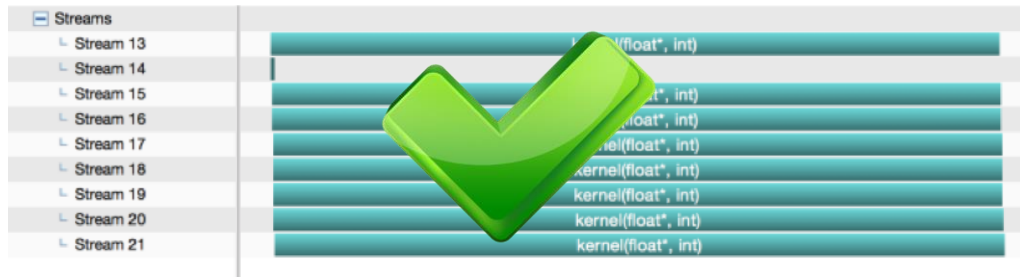



Dependency Scheduling Engine



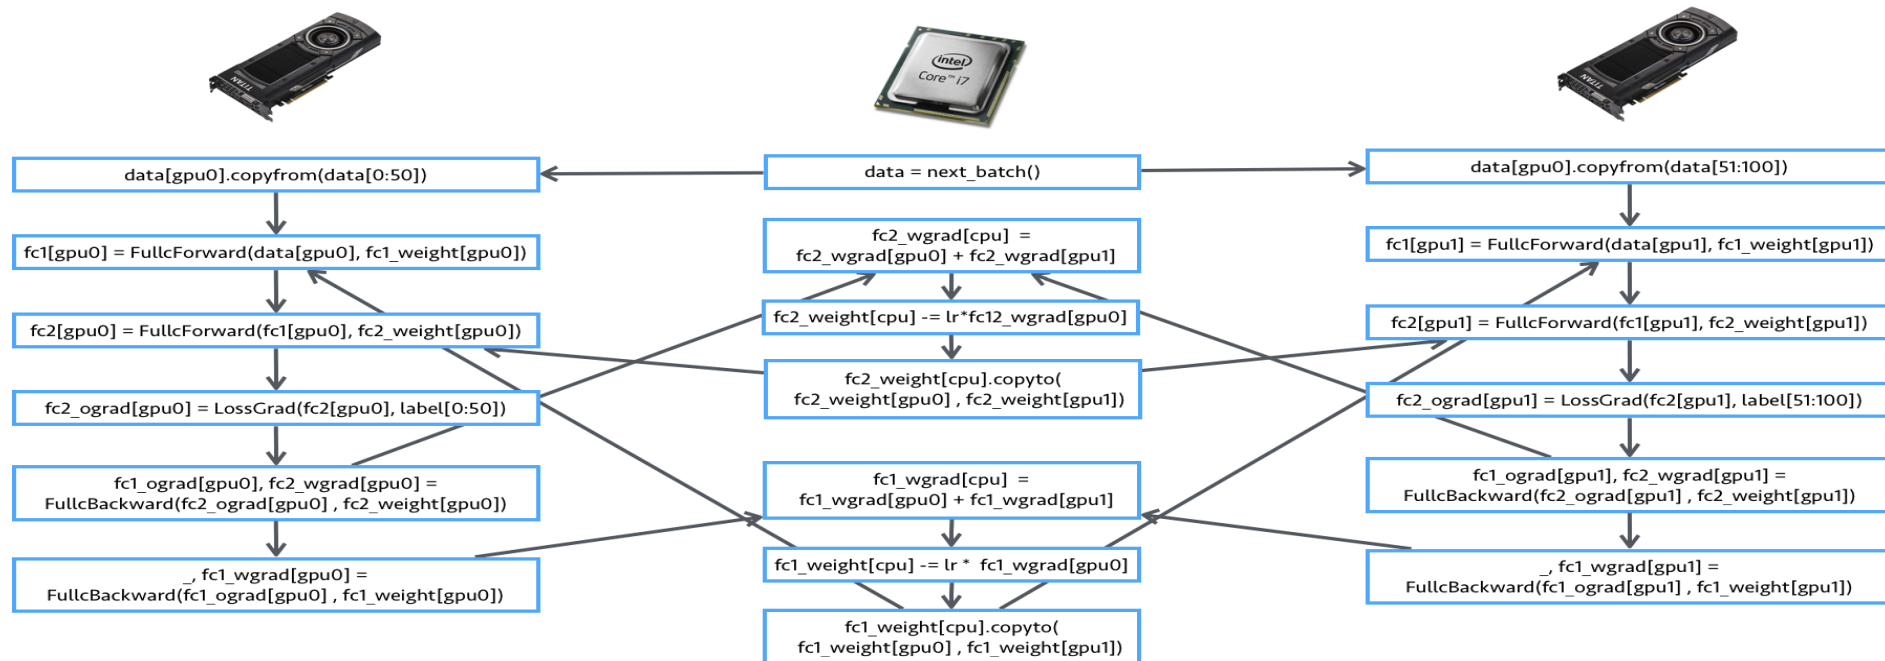
Need for Parallelism

- Speed is critical to deep learning
- Parallelism leads to higher performance
 - Parallelization across multiple GPUs
 - Parallel execution of small kernels
 - Overlapping memory transfer and computation
 - ...



Parallel Programs are Painful to Write...

- ... because of dependencies




Solution: Auto Parallelization with Dependency Engine

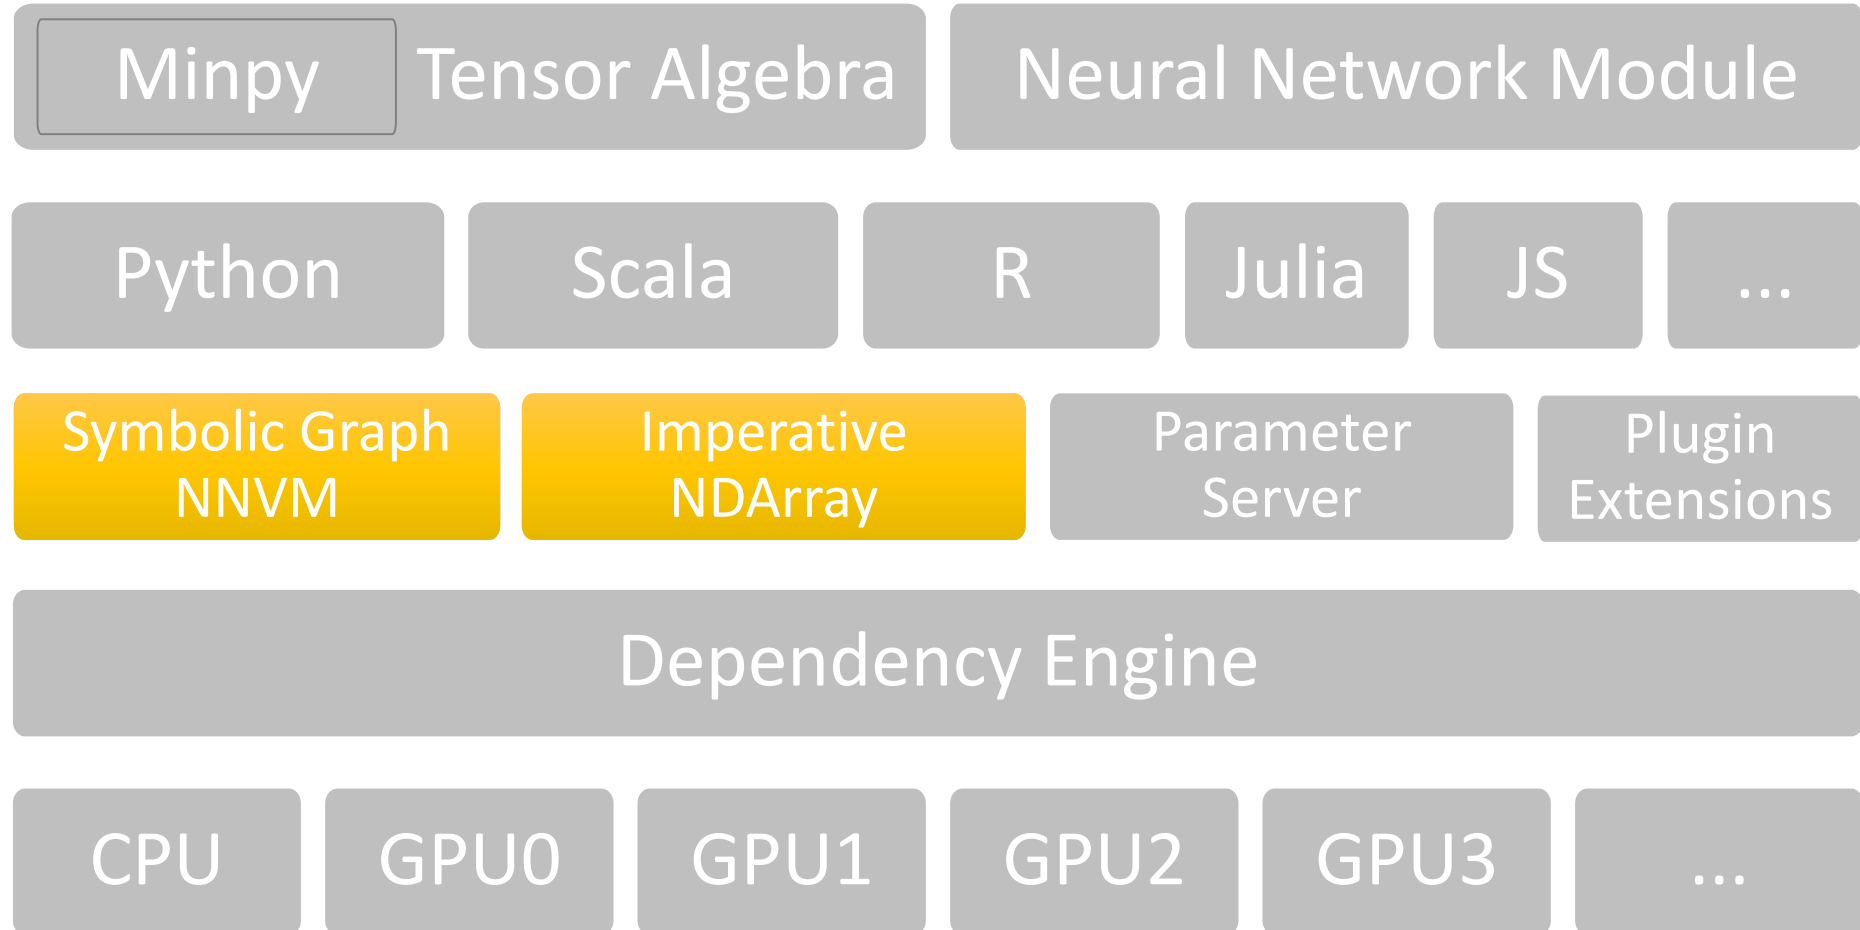
- Single thread abstraction of parallel environment

```
import mxnet as mx  
A = mx.nd.ones((2,2)) * 2  
C = A + 2  
B = A + 1  
D = B * C
```

Dependency
Engine

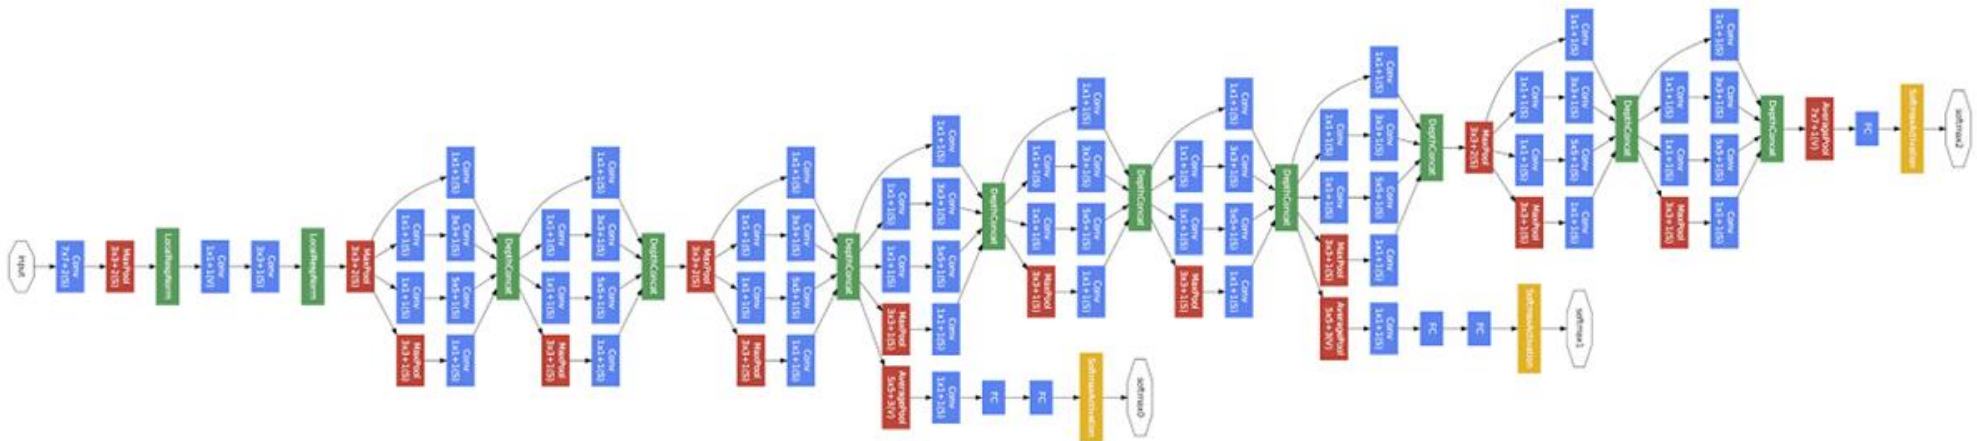


Symbolic vs. Imperative Deep Learning




Neural Networks as Symbolic Graph

- Most packages represent networks as graphs
 - MXNet, Caffe, Theano, Tensorflow, CNTK, ...
- Easy to store, port, and optimize
 - Parallelization, buffer sharing, operator fusion, serialize and deploy, ...



Deep Learning is More Than DAG

- All is good ...
 - ... until you start adding too many things to it ...



You just reinvented programming language

...

and lost most advantages of using graph


Neural Networks as Imperative Program

- DL, or ML in general, is largely tensor algebra.
 - Torch, Chainer, Matlab, R, Numpy, ...
- Imperative programs are flexible but hard to optimize:

```
import numpy as np
a = np.ones(10)
b = np.ones(10) * 2
c = b * a
print(c)
d = c + 1
```

Easy to tweak
with python
codes

```
A = Variable('A')
B = Variable('B')
C = B * A
D = C + 1
f = compile(D)
d = f(A=np.ones(10), B=np.ones(10)*2)
```




Neural Networks as Imperative Program

- DL, or ML in general, is largely tensor algebra.
 - Torch, Chainer, Matlab, R, Numpy, ...
- Imperative programs are flexible but hard to optimize:

```
import numpy as np
a = np.ones(10)
b = np.ones(10) * 2
c = b * a
d = c + 1
```

c cannot share memory with d,
because it could be used in future

```
A = Variable('A')
B = Variable('B')
C = B * A
D = C + 1
f = compile(D)
d = f(A=np.ones(10), B=np.ones(10)*2)
```



C can share memory with D,
because C cannot be seen by user

MXNet's Approach: Mixed Programming

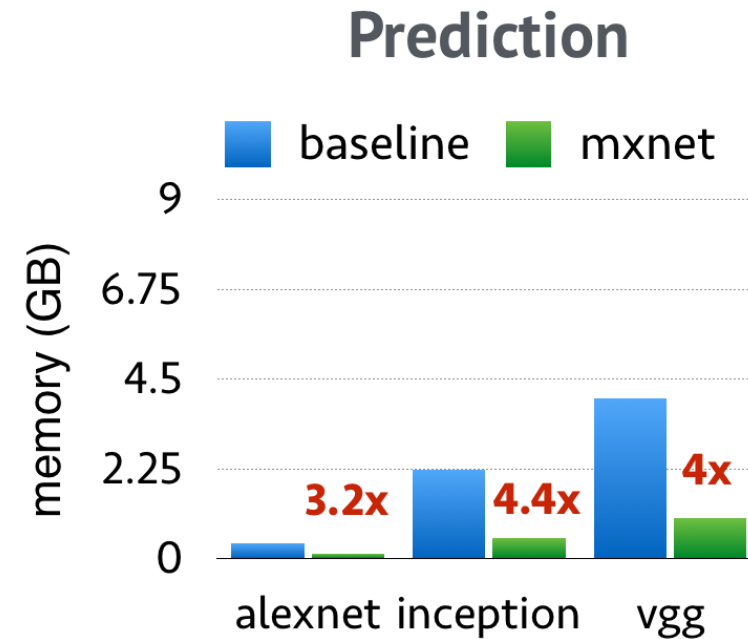

- Mix **symbolic** and **imperative** operations to get benefit of both.
- Symbolic graph: heavy, standard operations.
 - >90% of runtime, <10% of coding time.
- Imperative program: light, project specific operations.
 - <10% of runtime, >90% of coding time.
- Dependency Engine allows seamless combination of two parts.
- **Only optimize the bottleneck!**

MXNet's Approach: Static Graph

- In MXNet, graphs are **simple** and **static**:
 - Simple DAG without fancy logic
 - Fixed topology and tensor shapes
- Static graphs are:
 - **Faster to build**: multiple graphs instead of loop and condition.
 - **Easier to optimize**: Static memory planning, asynchronous execution, ...
- Flexibility compensated by imperative operations.

Smallest GPU Memory Footprint

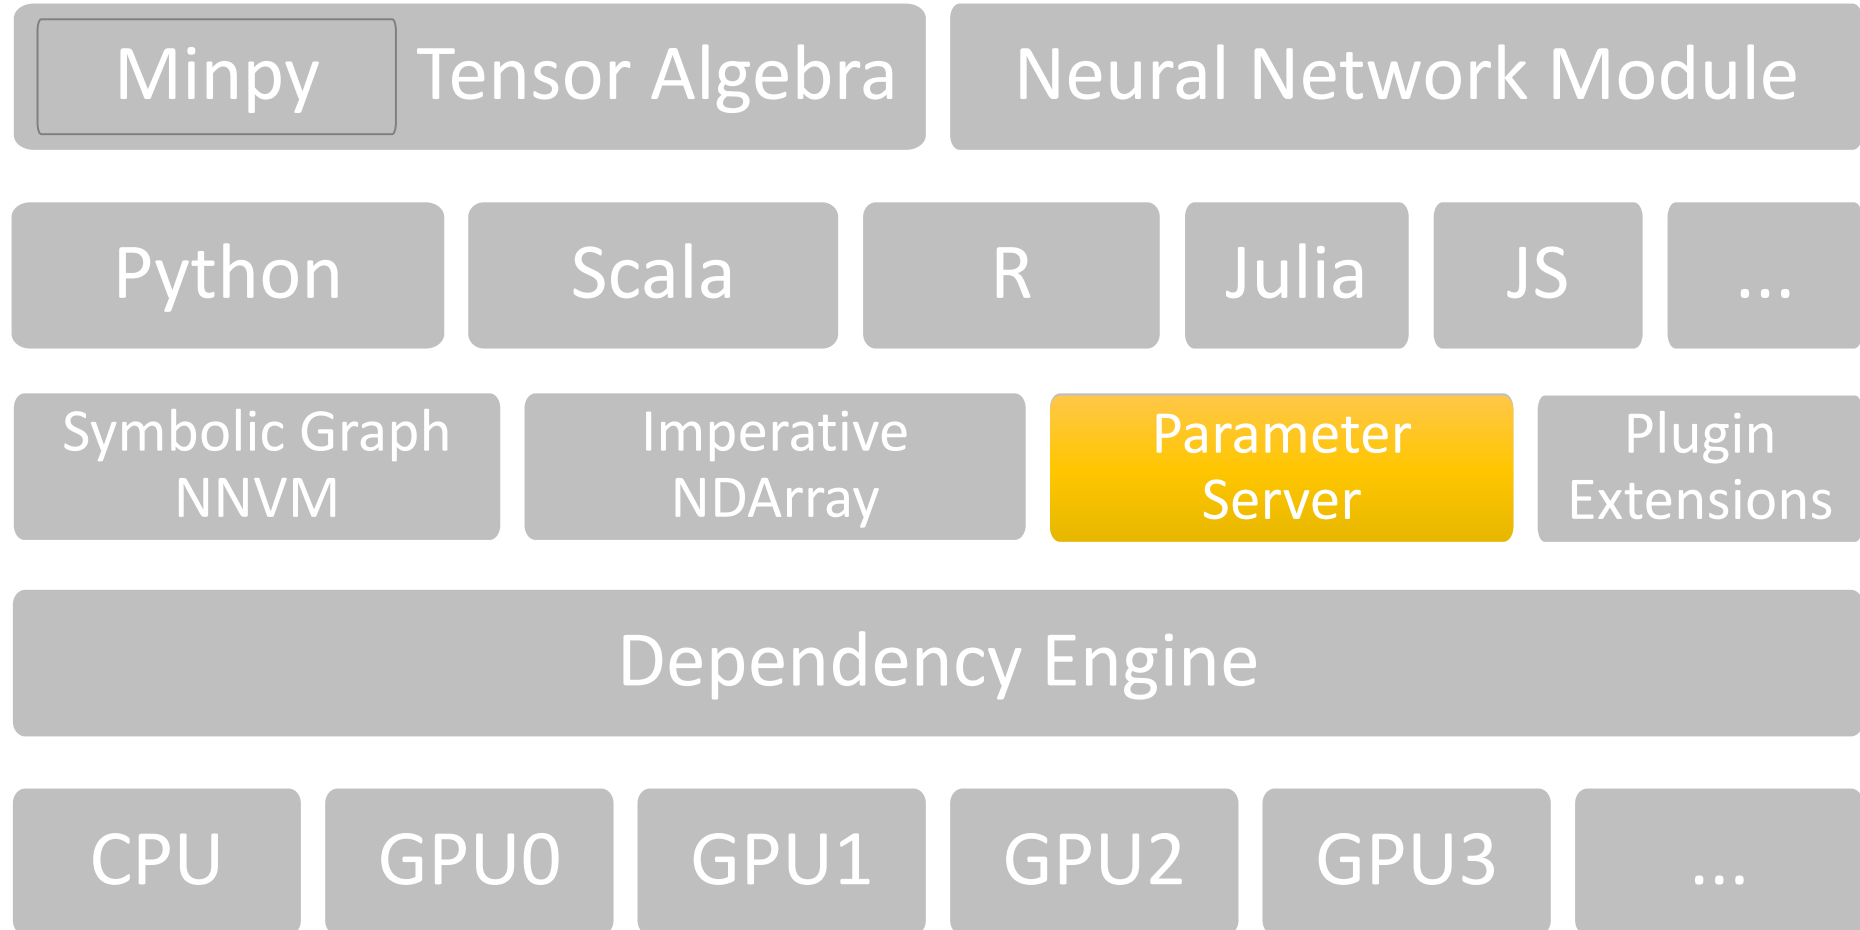
- Static graph enables aggressive memory sharing.



Trade Speed for Memory

- MXNet Mirror:
 - Discard result of small Ops on forward pass (ReLU, BN, etc).
 - Re-compute on backward pass.
 - 30%-50% memory saving at 90% speed.
- MXNet Memmonger:
 - Only keep result of \sqrt{N} (anchor layers) out of N layers on forward pass.
 - Re-compute \sqrt{N} layers between two anchor layers on backward pass.
 - Train $O(\sqrt{N})$ times larger model at 75% speed.

Parallel & Distributed Training



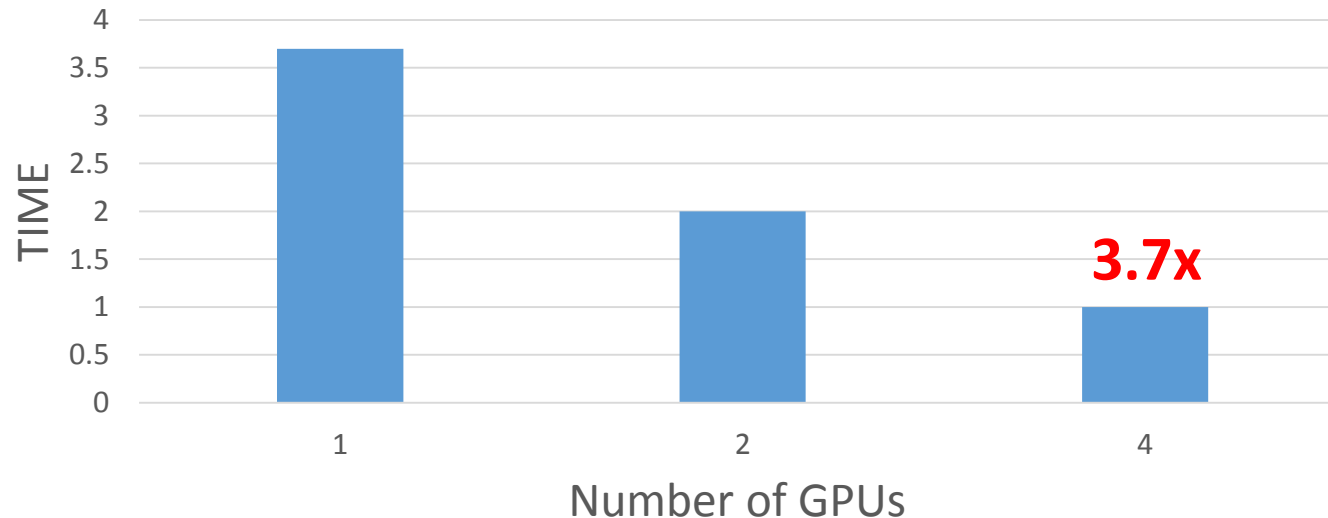
Drop-in Parallel Training

- Scaling to Multi-GPU machine as easy as one line change:

```
model = mx.mod.Module(net, ctx=mx.gpu(0))
```


```
-> model = mx.mod.Module(net, ctx=[mx.gpu(0), mx.gpu(1)])
```

- Near linear speedup on a single machine:



Parallel Training: Under the Hood

- Read a data partition
- Pull the weight
- Compute the gradient
- Push the gradient
- Update the weight



Distributed Training

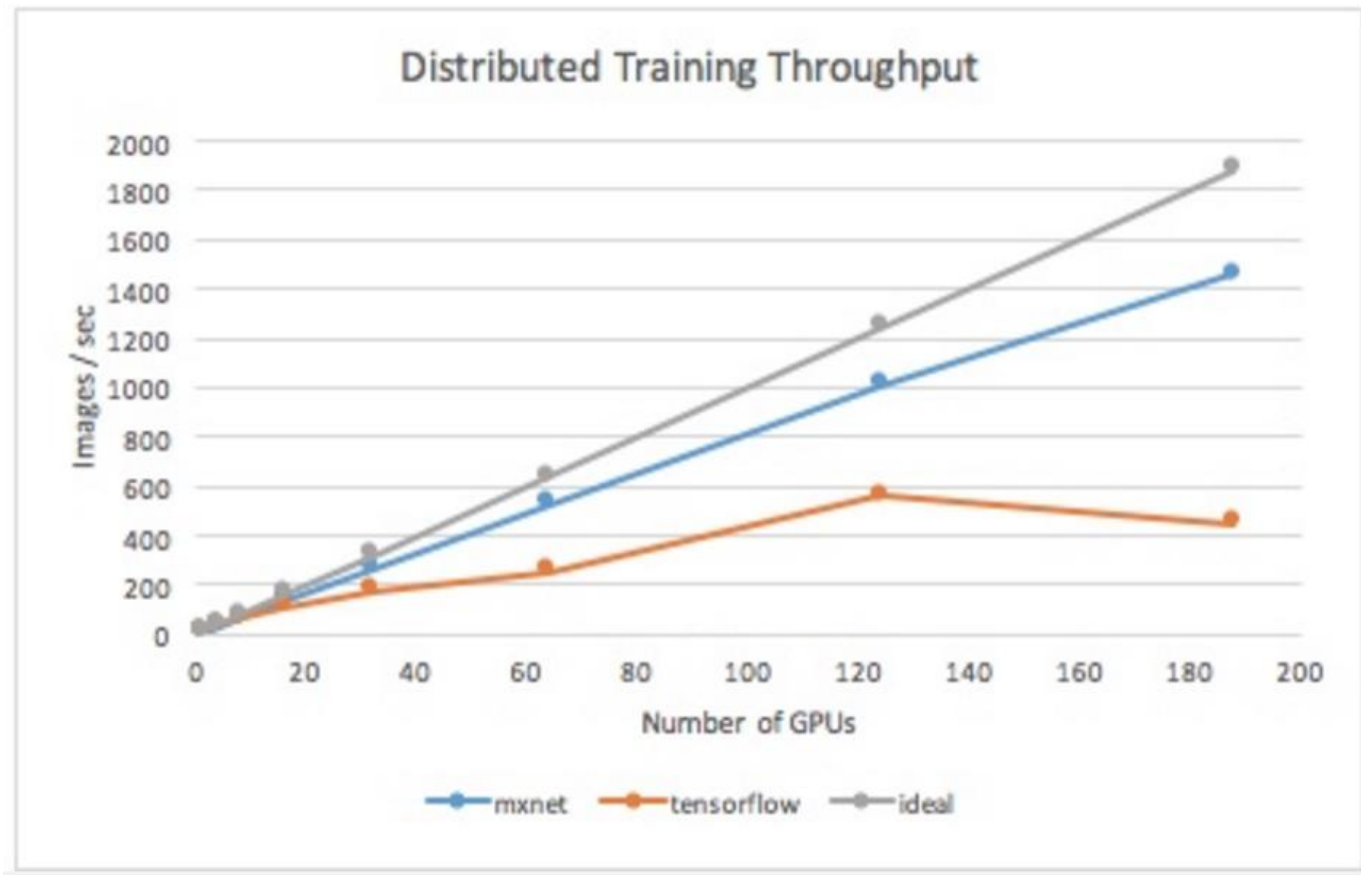
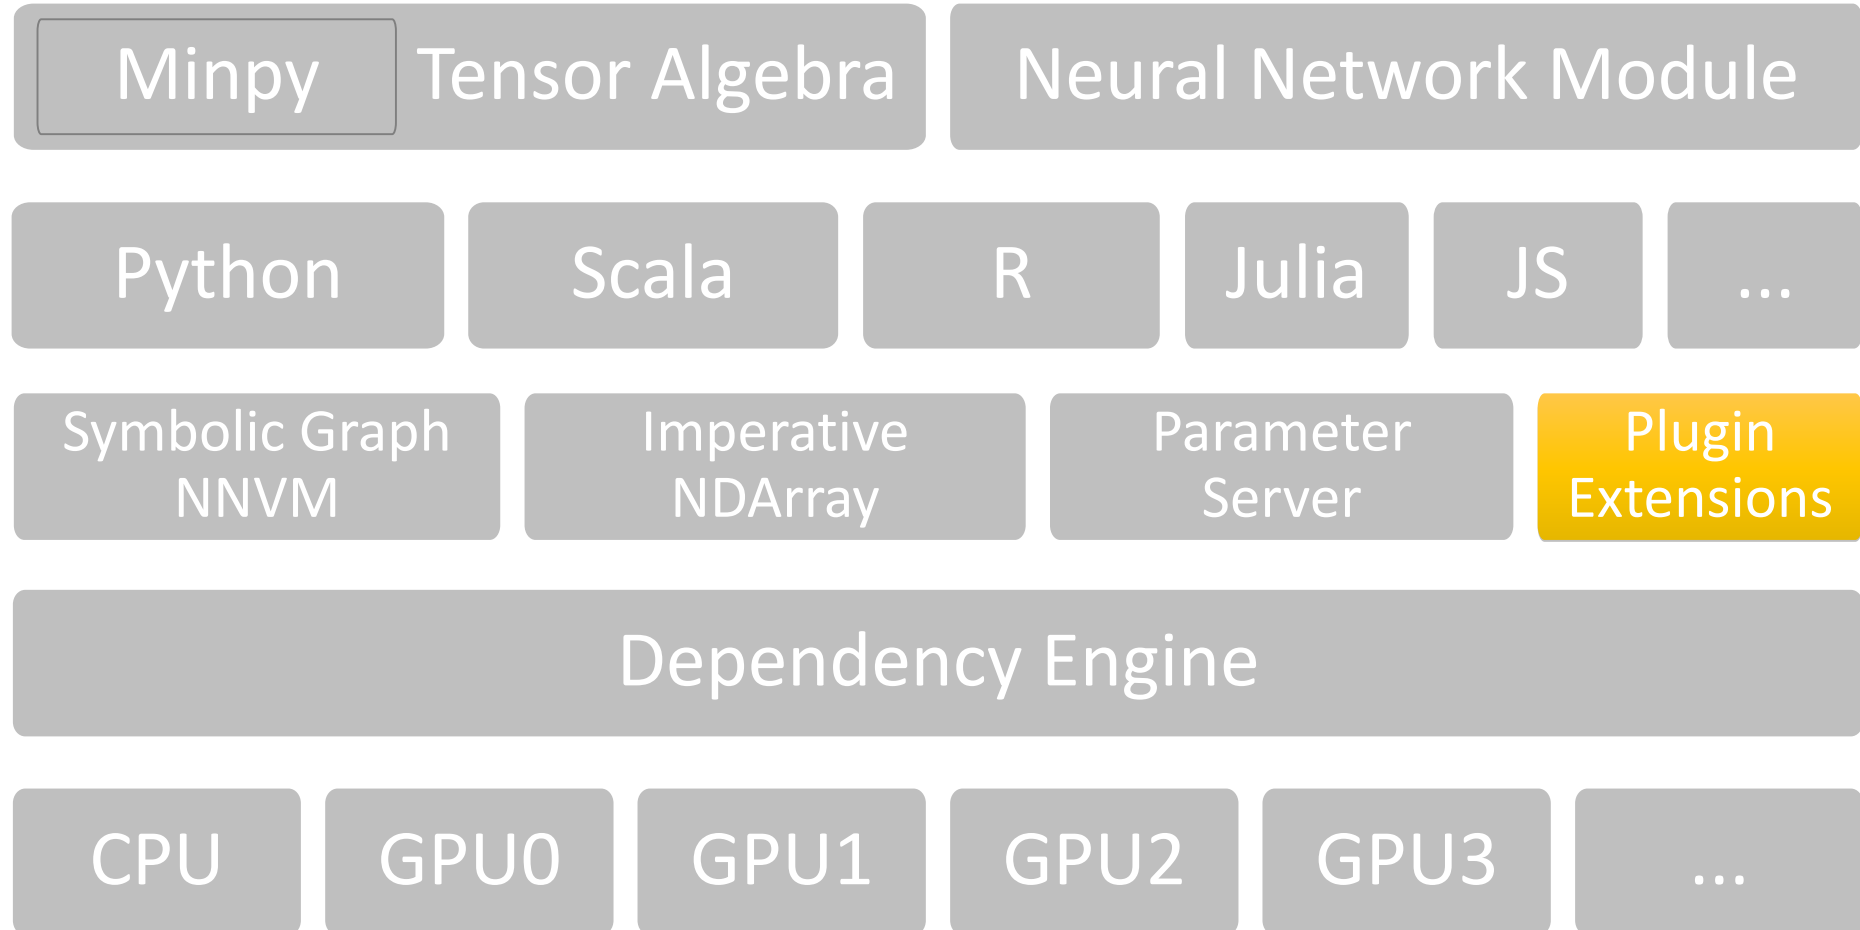


Figure by Carlos Guestrin @Turi

- Scale to multi-machine with the same key-value store interface.
- 2x faster than Tensorflow on >10 4GPU machines
- Latest Update: 74x acceleration on 10 machines with 8GPUs in each machine.

Parallel & Distributed Training



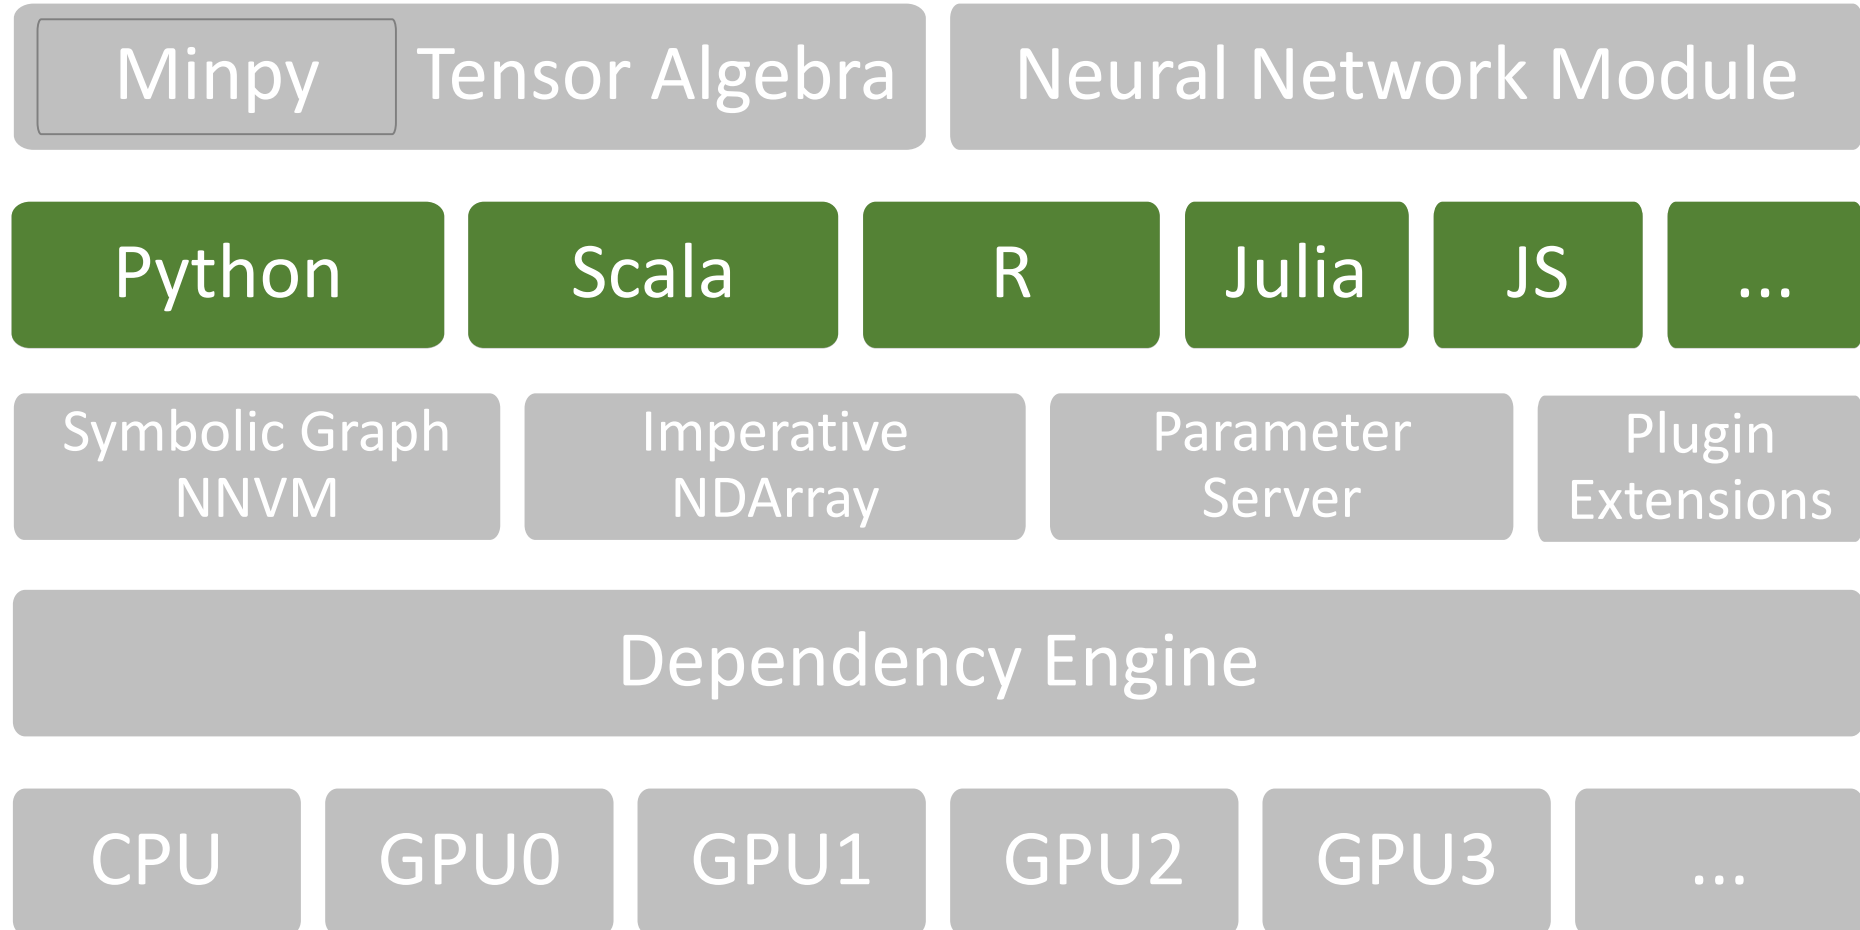
Plugin Extensions

- TorchModule:
 - Use Torch NN layers and tensor functions in MXNet graph.
`fc1 = mx.sym.TorchModule(lua_string='nn.Linear(784, 128)', ...)`
- CaffeOp:
 - Use Caffe layers in MXNet graph.
`fc1 = mx.symbol.CaffeOp(prototxt="layer{type:\"InnerProduct\" inner_product_param{num_output: 128} }", ...)`
- WrapCTC:
 - Use Baidu's CTC module for sequence learning in MXNet.
- OpenCV:
 - Multi-threaded OpenCV interface that by pass GIL for fast image IO.


Mainstream Applications in Vision/NLP/Speech

- Image Classification
 - Inception/ResNet
- Object Detection
 - Faster RCNN
- Image Segmentation
 - FCN/Deeplab
- OCR
 - Warp-CTC
- Char LSTM/Char CNN/Speech Acoustic Modeling/Neural Art...


Runs Everywhere



Code with Any Language



Lines of code



Train in the Cloud

Load data from
distributed filesystems



HDFS



S3



Blob

⋮

multithreaded read/write
to hide network latency

Launch distributed jobs




SSH



MPI



qsub



Yarn



⋮


easily extend to other cluster
resource management software

Deploy Everywhere

Beyond   

Amalgamation

- ✦ Fit the core library with all dependencies into a single C++ source file
- ✦ Easy to compile on   ...

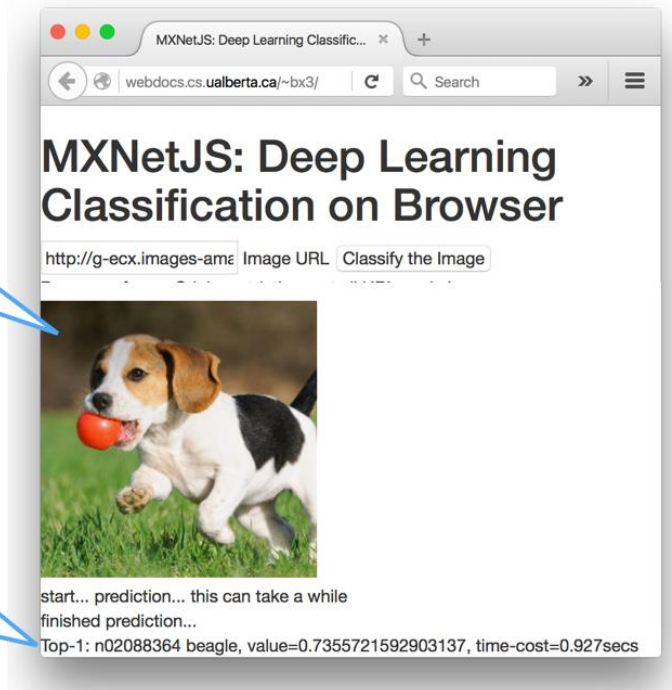


BlindTool by Joseph Paul Cohen, demo on Nexus 4

The first image for search "dog" at images.google.com

Outputs "beagle"
with prob = 73%
within 1 sec

Runs in browser
with Javascript



Thanks